

Cutting our p99 in half

Cutting our p99 in half

Ben Armstrong · Platform team

Where the time went

- Serialising on a single writer
- Retries with no jitter
- A cache that never warmed



db · queue · cache

Cutting our p99 in half

One writer, many waiters

Every request queued behind the same lock, so throughput was flat no matter how many replicas we added.

What we changed

1. Sharded the writer by tenant
2. Added jitter to every retry path
3. Warmed the cache on deploy

Cutting our p99 in half

YouTube/GIF video support



an embedded GIF

► <https://www.youtube.com/watch?v=YE7VzlLtp-4>

a YouTube link

Cutting our p99 in half

The result

840ms → 390ms

p99, measured over four weeks

Cutting our p99 in half

Questions